

```
In [1]: # Created by: Michael Cullen
# 19/11/2024
# https://zenodo.org/records/10457828
# https://www.kaggle.com/datasets/syedanwarafri/vehicle-sales-data/data
```

```
In [2]: import pandas as pd
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from sklearn.metrics import mean_squared_error, r2_score
```

```
In [3]: # made by chatgpt to install csv as file is too big to commit to github

import kagglehub
import shutil
import os

# Define the current working directory
current_directory = os.getcwd()

# Check if the CSV file already exists in the current directory
csv_exists = any(file.endswith(".csv") for file in os.listdir(current_directory))

if not csv_exists:
    # Download the dataset using kagglehub
    default_path = kagglehub.dataset_download("syedanwarafri/vehicle-sales-data")

    # Move all downloaded CSV files to the current directory
    for file_name in os.listdir(default_path):
        if file_name.endswith(".csv"):
            shutil.move(os.path.join(default_path, file_name), os.path.join(current_directory, file_name))

    # Delete the downloaded folder after moving the files
    shutil.rmtree(default_path)

    print("Dataset files moved to:", current_directory)
    print(f"Deleted temporary folder: {default_path}")
else:
    print("CSV file already exists in the current directory.")
```

CSV file already exists in the current directory.

```
In [4]: df = pd.read_csv('car_prices.csv')

df.head()
```

Out[4]:

	year	make	model	trim	body	transmission	vin	state	condition
0	2015	Kia	Sorento	LX	SUV	automatic	5xyktca69fg566472	ca	5.0
1	2015	Kia	Sorento	LX	SUV	automatic	5xyktca69fg561319	ca	5.0
2	2014	BMW	3 Series	328i SULEV	Sedan	automatic	wba3c1c51ek116351	ca	45.0
3	2015	Volvo	S60	T5	Sedan	automatic	yv1612tb4f1310987	ca	41.0
4	2014	BMW	6 Series Gran Coupe	650i	Sedan	automatic	wba6b2c57ed129731	ca	43.0

In [5]:

```
print("Amount of columns: ", df.shape[0])
```

Amount of columns: 558837

In [6]:

```
df.isnull().sum()
```

Out[6]:

year	0
make	10301
model	10399
trim	10651
body	13195
transmission	65352
vin	4
state	0
condition	11820
odometer	94
color	749
interior	749
seller	0
mmr	38
sellingprice	12
saledate	12

dtype: int64

```
In [7]: # Drop not needed columns
df = df.drop(columns=['trim','vin','state','interior','seller','saledate'], errors=

# Remove rows with missing data
df = df.dropna()

df.head(20)
```

```
Out[7]:
```

	year	make	model	body	transmission	condition	odometer	color	mm
0	2015	Kia	Sorento	SUV	automatic	5.0	16639.0	white	20500.0
1	2015	Kia	Sorento	SUV	automatic	5.0	9393.0	white	20800.0
2	2014	BMW	3 Series	Sedan	automatic	45.0	1331.0	gray	31900.0
3	2015	Volvo	S60	Sedan	automatic	41.0	14282.0	white	27500.0
4	2014	BMW	6 Series Gran Coupe	Sedan	automatic	43.0	2641.0	gray	66000.0
5	2015	Nissan	Altima	Sedan	automatic	1.0	5554.0	gray	15350.0
6	2014	BMW	M5	Sedan	automatic	34.0	14943.0	black	69000.0
7	2014	Chevrolet	Cruze	Sedan	automatic	2.0	28617.0	black	11900.0
8	2014	Audi	A4	Sedan	automatic	42.0	9557.0	white	32100.0
9	2014	Chevrolet	Camaro	Convertible	automatic	3.0	4809.0	red	26300.0
10	2014	Audi	A6	Sedan	automatic	48.0	14414.0	black	47300.0
11	2015	Kia	Optima	Sedan	automatic	48.0	2034.0	red	15150.0
12	2015	Ford	Fusion	Sedan	automatic	2.0	5559.0	white	15350.0
13	2015	Kia	Sorento	SUV	automatic	5.0	14634.0	silver	20600.0
15	2015	Nissan	Altima	Sedan	automatic	2.0	11398.0	black	14750.0
17	2014	Audi	Q5	SUV	automatic	49.0	7983.0	white	37100.0
18	2014	Chevrolet	Camaro	Coupe	automatic	17.0	13441.0	black	17750.0
19	2014	BMW	6 Series	Convertible	automatic	34.0	8819.0	black	68000.0
20	2015	Chevrolet	Impala	Sedan	automatic	19.0	14538.0	silver	24300.0
21	2014	BMW	5 Series	Sedan	automatic	29.0	25969.0	black	34200.0

```
In [8]: print(df.isnull().sum())
print("Amount of columns: ", df.shape[0])
```

```
year            0
make            0
model           0
body            0
transmission    0
condition       0
odometer        0
color           0
mmr             0
sellingprice    0
dtype: int64
Amount of columns: 472325
```

```
In [9]: from sklearn.preprocessing import LabelEncoder

# Convert odometer, year, and mmr to numeric
df['odometer'] = pd.to_numeric(df['odometer'], errors='coerce')
df['year'] = pd.to_numeric(df['year'], errors='coerce')
df['mmr'] = pd.to_numeric(df['mmr'], errors='coerce')

# Define the function to categorise odometer values
def categorise_odometer(odometer): # 1 = 0-20k, 2 = 20-50k, 3 = 50-100k, 4 = 100-1
    if odometer <= 20000:
        return 1 # 0-20k
    elif odometer <= 50000:
        return 2 # 20-50k
    elif odometer <= 100000:
        return 3 # 50-100k
    elif odometer <= 150000:
        return 4 # 100-150k
    elif odometer <= 200000:
        return 5 # 150-200k
    else:
        return 6 # 200k+

# Define function to categorise year
def categorise_year(year): # 1 = very old, 2 = old, 3 = modern, 4 = new
    if year < 1990:
        return 1 # Very old
    elif year < 2005:
        return 2 # Old
    elif year < 2015:
        return 3 # Modern
    else:
        return 4 # New

# Define function to categorise price
def categorise_price(price): # 1 = Low, 2 = medium, 3 = high, 4 = premium
    if price < 5000:
        return 1 # Low
    elif price < 20000:
        return 2 # Medium
    elif price < 40000:
        return 3 # High
    else:
        return 4 # Premium
```

```
# Define function to categorise mmr
def categorise_mmr(mmr): # 1 = very low, 2 = low, 3 = medium, 4 = high, 5 = very h
    if mmr < 5000:
        return 1 # Very Low
    elif mmr < 15000:
        return 2 # Low
    elif mmr < 30000:
        return 3 # Medium
    elif mmr < 50000:
        return 4 # High
    else:
        return 5 # Very High

# Apply categorisation functions
df['odometer_group'] = df['odometer'].apply(categorise_odometer)
df['year_group'] = df['year'].apply(categorise_year)
df['price_group'] = df['sellingprice'].apply(categorise_price)
df['mmr_group'] = df['mmr'].apply(categorise_mmr)

# Group rare colors into 'Other' to reduce noise
if 'color' in df.columns:
    color_counts = df['color'].value_counts()
    rare_colors = color_counts[color_counts < 50].index # Adjust threshold as need
    df['color'] = df['color'].replace(rare_colors, 'Other')

# Handle categorical features using LabelEncoder
categorical_columns = ['make', 'model', 'body', 'cylinders', 'drive', 'condition',

# Encode all categorical columns
for col in categorical_columns:
    if col in df.columns: # Ensure the column exists in the dataset
        df[col] = LabelEncoder().fit_transform(df[col].astype(str))

# Drop the original columns
df = df.drop(columns=['odometer', 'year', 'sellingprice', 'mmr'], errors='ignore')

# Show the first 10 rows of the processed DataFrame
df.head(10)
```

```
Out[9]:
```

	make	model	body	transmission	condition	color	odometer_group	year_group	price
0	24	637	35		0	40	16	1	4
1	24	637	35		0	40	16	1	4
2	3	8	36		0	35	8	1	3
3	51	575	36		0	31	16	1	4
4	3	33	36		0	33	8	1	3
5	36	63	36		0	0	8	1	4
6	3	406	36		0	24	2	1	3
7	7	186	36		0	10	2	2	3
8	2	46	36		0	32	16	1	3
9	7	135	9		0	20	13	1	3

```
In [10]: # Define features (X) and target (y)
X = df.drop(columns='price_group')
y = df['price_group']

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_sta
```

```
In [11]: from sklearn.model_selection import GridSearchCV
#https://www.kaggle.com/code/syedanwarafриди/eda-visualizaton-random-forest-on-vehi
# code to find out best parameters sourced from kaggle forum

# Random Forest hyperparameter tuning using GridSearchCV
rf_param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [30, 40, 50],
    'min_samples_split': [8, 12],
    'min_samples_leaf': [4, 5]
}


# Instantiate the model using the best parameters
rf_model = RandomForestRegressor(random_state=42)
rf_best = GridSearchCV(rf_model, rf_param_grid, cv=5, scoring='neg_mean_squared_err

# Fit the model to the training data
rf_best.fit(X_train, y_train)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

```
c:\Users\mjcul\OneDrive\Documents\GitHub\DataScience\venv\lib\site-packages\numpy\ma
\core.py:2881: RuntimeWarning: invalid value encountered in cast
    _data = np.array(data, dtype=dtype, copy=copy,
```

Out[11]:



```
In [12]: # Make predictions
y_pred = rf_best.predict(X_test)


# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse:.2f}')

r2 = r2_score(y_test, y_pred)
print(f'R-squared: {r2:.2f}')
```

Mean Squared Error: 0.08

R-squared: 0.82

```
In [14]: importances = rf_best.best_estimator_.feature_importances_
plt.barh(X.columns, importances)
plt.xlabel("Feature Importance")
plt.title("Random Forest Feature Importances")
plt.show()
```



```
In [ ]: # Scatter plot for actual vs predicted values
plt.figure(figsize=(10, 6))
```

```
plt.scatter(y_test, y_pred, alpha=0.5)
plt.title('Actual vs Predicted Values')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.grid(True)
plt.show()
```

